Time Series Study Sheet

Sarra Lajnef
Based on: Applied Time Series Analysis (Terence C. Mills)

April 13, 2026

1 Part 1: Time Series and Their Features

1.1 Fundamental Definitions

A **time series** on a variable x is denoted as x_t, where t is the time index.

- **Observation Period:** $t = 1, 2, \ldots, T$, where T is the last observation.
- **Forecast Horizon:** Future values are denoted as x_{T+h}, where h is the horizon.
- **Sequential Order:** Unlike cross-sectional data, the order of observations in a time series is paramount because of potential temporal dependence.

1.2 Key Features of Time Series

1.2.1 Autocorrelation and Periodic Movements

Autocorrelation refers to the correlation of a series with its own past values. Many series exhibit periodic movements that are not strictly seasonal but repeat over certain intervals.

1.2.2 Seasonality

Seasonality consists of regular, short-term cycles that repeat within a fixed period, typically a year.

1.2.3 Stationarity and Nonstationarity

A process is **stationary** if its statistical properties (mean, variance) remain constant over time.

- **Nonstationarity:** Common in real-world data where the mean or variance changes over time.
- **Integration:** A series is integrated of order d, denoted $I(d)$, if it must be differenced d times to become stationary.
- **First-Difference Operator (Δ):** Used to transform a nonstationary series:

$$\Delta x_t = x_t - x_{t-1} \tag{1}$$

1.2.4 Trends

A trend represents the long-term evolution of a series.

- **Deterministic Trend:** Often modeled as a linear function of time: $x_t = \beta_0 + \beta_1 t + \epsilon_t$.
- **Stochastic Trend:** Represented by a "random walk" (explained afterwards) where the trend evolves based on cumulative random shocks.

1.2.5 The Concept of the "Level" of a Series (for me didn't do it in class

The **level** of a time series represents its local mean or the underlying value at a specific point in time, once noise and seasonal fluctuations are removed.

- **In Stationary Series:** The level is constant and corresponds to the long-run mean μ.
- **In Non-stationary Series:** The level evolves over time.
  In a Random Walk ($x_t = x_{t-1} + \epsilon_t$), the current level is the sum of all past shocks.
- **State-Space Representation:** In many decomposition models (like Exponential Smoothing), the level L_t is updated recursively:

$$L_t = \alpha x_t + (1 - \alpha)(L_{t-1} + T_{t-1}) \tag{2}$$

  where $0 < \alpha < 1$ is the smoothing parameter, x_t is the current observation, and T_{t-1} is the previous trend estimate.
- **Level Shift:** A structural break where the mean of the series undergoes a permanent jump or drop. This is often modeled using a dummy variable D_t:

$$x_t = \mu + \delta D_t + \epsilon_t, \quad D_t = \begin{cases} 0 & t < \tau \\ 1 & t \geq \tau \end{cases} \tag{3}$$

  where τ is the point of the break and δ is the magnitude of the shift.

1.2.6 Volatility

Volatility refers to a time-varying variance.

- **Volatility Clustering:** Periods of high variation followed by more high variation, common in financial data.
- **Model:** Often handled using GARCH (Generalized Autoregressive Conditional Heteroskedastic) models.

1.3 Natural Constraints

Some series must follow mathematical constraints due to their nature:

- **Count Data:** Observations are non-negative integers $(0, 1, 2 \dots)$.
- **Compositional Data:** Series representing shares of a whole, satisfying:

$$0 \leq x_{i,t} \leq 1 \quad \text{and} \quad \sum_{i=1}^{n} x_{i,t} = 1 \tag{4}$$

1.4 Multivariate Features

- **Common Trends:** Two or more nonstationary series may be "bound together" by a common underlying trend (Cointegration).
- **Causality:** One series may help predict another (Granger Causality).

2 Part 2: Transforming Time Series

Prior to statistical analysis, it is often necessary to transform data to make it more amenable to modeling. There are three general classes of transformations: distributional, stationarity-inducing, and decompositional.

2.1 Distributional Transformations

These aim to achieve approximate normality, symmetry, and constant variance.

2.1.1 Logarithmic Transformation

A popular choice for non-negative series that are skewed to the right. It "straightens out" exponential trends and stabilizes variance.

- **Usage:** Effective when the standard deviation is proportional to the level of the series.
- **Formula:** $y_t = \log(x_t)$.

Figure 1: serie skewed to the right

2.1.2 Box-Cox Transformation

A general class of power transformations that includes the logarithm as a special case for positive x:

$$f_{BC}(x_t; \lambda) = \begin{cases} (x_t^\lambda - 1)/\lambda & \lambda \neq 0 \\ \log x_t & \lambda = 0 \end{cases} \tag{5}$$

2.1.3 Handling Negative Values

For series with negative values, several alternatives exist:

- **Signed Power:** $f_{SP}(x_t; \lambda) = \text{sgn}(x_t)|x_t|^\lambda - 1/\lambda$ for $\lambda > 0$.
- **Inverse Hyperbolic Sine (IHS):** Useful for extreme values of either sign:

$$f_{IHS}(x_t; \lambda) = \frac{\sinh^{-1}(\lambda x_t)}{\lambda} = \frac{\log(\lambda x_t + (\lambda^2 x_t^2 + 1)^{1/2})}{\lambda} \tag{6}$$

2.2 Stationarity Inducing Transformations

To remove deterministic or stochastic trends, differencing is applied.

2.2.1 First Differencing

Used for series that are integrated of order 1 ($I(1)$):

$$\Delta x_t = (1 - B)x_t = x_t - x_{t-1} \tag{7}$$

where B is the backshift operator ($Bx_t = x_{t-1}$).

2.2.2 Seasonal Differencing

For a series with a seasonal cycle of length s (e.g., $s = 12$ for monthly data):

$$\Delta_s x_t = (1 - B^s)x_t = x_t - x_{t-s} \tag{8}$$

2.3 Decomposing a Time Series

Observed series X_t are often decomposed into Trend (T_t), Seasonal (S_t), and Irregular (I_t) components.

2.3.1 Additive vs. Multiplicative Models

- **Additive:** $X_t = T_t + S_t + I_t$.
- **Multiplicative:** $X_t = T_t \times S_t \times I_t$ (often transformed to additive via logs).

2.3.2 Smoothing Transformations (Moving Averages)

A simple way to extract the trend is via a symmetric $2k + 1$ point moving average:

$$M_t = \frac{1}{2k+1} \sum_{j=-k}^{k} x_{t+j} \tag{9}$$

2.3.3 Hodrick-Prescott (HP) Filter

A more advanced decomposition that minimizes the following sum for a chosen smoothing parameter λ:

$$\min \sum_{t=1}^{T} (x_t - T_t)^2 + \lambda \sum_{t=2}^{T-1} [(T_{t+1} - T_t) - (T_t - T_{t-1})]^2 \tag{10}$$

Common values for λ include 14,400 for monthly data or 1,600 for quarterly data.

3 Part 3: ARMA Models for Stationary Time Series

3.1 Intuitive Interpretation of AR and MA Components

Before diving into the complex estimations, it is helpful to understand the behavioral "personality" of these models.

3.1.1 AR (Autoregressive): The Long-Term Memory

The **AR** component describes a system that depends on its **own previous values**.

- **Analogy:** It is like a person's mood; if someone was happy yesterday, he is likely to still be somewhat happy today.
- **Mechanism:** Each new value is a weighted percentage of the last value, plus a current "random surprise."
- **Impact:** It creates **momentum** or **persistence**. For example, if the temperature is 30°C today, it is unlikely to drop to 0°C tomorrow because the "memory" of today's heat lingers.

3.1.2 MA (Moving Average): The Shock Absorber

The **MA** component does **not** depend on the previous values of the series, but rather on the **previous random shocks** (surprises or errors).

- **Analogy:** Imagine an earthquake. The "shock" happens today. Tomorrow, we feel an "aftershock." By the third day, that specific shock is gone.
- **Mechanism:** It looks at the "random noise" from yesterday and determines how much of that noise should still ripple into today.
- **Impact:** It handles **short-term echoes**. Once the shock exceeds the order q, it is completely forgotten.

3.1.3 ARMA: The Combined Personality

Combining them into an **ARMA** model allows us to represent a series with:

- **Persistence (AR):** The tendency to stay where it was.
- **Reaction to Shocks (MA):** The tendency to ripple for a short time after a sudden surprise.

3.1.4 Summary Comparison

Feature	**AR (Autoregressive)**	**MA (Moving Average)**
Logic	"I am where I am because of where I was."	"I am where I am because of the 'news' I just received."
Memory	Long-term (fades geometrically).	Short-term (cuts off abruptly).
Analogy	**Momentum:** Like a car rolling after we stop hitting the gas.	**Echo:** Like a sound bouncing for a moment, then vanishing.

3.2 Theoretical Foundations

3.2.1 The Stochastic Process

A **time series** is a realization of a **stochastic process** $\{x_t, t \in \mathbb{Z}\}$. We assume the process is **linearly regular**, meaning it can be expressed as a linear combination of its past innovations.

3.2.2 Stationarity (Covariance Stationarity)

A process is weakly stationary if:

1. $E[x_t] = \mu$ (The mean is constant for all t).
2. $Var(x_t) = E[(x_t - \mu)^2] = \sigma^2 = \gamma_0 < \infty$.
3. $Cov(x_t, x_{t-k}) = E[(x_t - \mu)(x_{t-k} - \mu)] = \gamma_k$ for all t.

The **Autocorrelation Function (ACF)** is defined as:

$$\rho_k = \frac{\gamma_k}{\gamma_0} \tag{11}$$

Note that $\rho_0 = 1$ and $|\rho_k| \leq 1$.

3.2.3 The White Noise Process

The building block of all ARMA models is **White Noise** $\{a_t\}$, defined as:

- $E[a_t] = 0$
- $E[a_t^2] = \sigma_a^2$
- $E[a_t a_s] = 0$ for $t \neq s$ (uncorrelated shocks).

3.3 Autoregressive Processes: $AR(p)$

3.3.1 The $AR(1)$ Model: Detailed Derivation

The model is $x_t = \phi x_{t-1} + a_t$.
By repeated substitution (assuming the process started in the infinite past):

$$x_t = a_t + \phi a_{t-1} + \phi^2 a_{t-2} + \cdots = \sum_{j=0}^{\infty} \phi^j a_{t-j} \tag{12}$$

Variance calculation:

$$\gamma_0 = Var\left(\sum_{j=0}^{\infty} \phi^j a_{t-j}\right) = \sigma_a^2 \sum_{j=0}^{\infty} \phi^{2j} = \frac{\sigma_a^2}{1-\phi^2} \tag{13}$$

This requires the **Stationarity Condition**: $|\phi| < 1$.

3.3.2 The General $AR(p)$ and Yule-Walker Equations

The $AR(p)$ model is $x_t = \phi_1 x_{t-1} + \cdots + \phi_p x_{t-p} + a_t$.
To find the autocorrelations, we multiply by x_{t-k} and take expectations. This yields the **Yule-Walker Equations**:

$$\rho_k = \phi_1 \rho_{k-1} + \phi_2 \rho_{k-2} + \cdots + \phi_p \rho_{k-p} \quad \text{for } k > 0 \tag{14}$$

This system of equations allows us to solve for the parameters ϕ_i if we know the autocorrelations ρ_k.

3.4 Moving Average Processes: $MA(q)$

3.4.1 The $MA(1)$ Model

Defined as $x_t = a_t - \theta a_{t-1}$.

- **Mean:** $E[x_t] = E[a_t] - \theta E[a_{t-1}] = 0$.
- **Variance:** $\gamma_0 = E[(a_t - \theta a_{t-1})^2] = \sigma_a^2(1+\theta^2)$.
- **Autocovariance at lag 1:** $\gamma_1 = E[(a_t - \theta a_{t-1})(a_{t-1} - \theta a_{t-2})] = -\theta\sigma_a^2$.
- **ACF:** $\rho_1 = \frac{-\theta}{1+\theta^2}$ and $\rho_k = 0$ for $k > 1$.

3.4.2 Invertibility

An $MA(1)$ is **invertible** if $|\theta| < 1$.
If invertible, we can express the error a_t as an infinite sum of past observations:

$$a_t = \sum_{j=0}^{\infty} \theta^j x_{t-j} \tag{15}$$

3.4.3 General $MA(q)$ Autocorrelation

For a general $MA(q)$ model $x_t = a_t - \theta_1 a_{t-1} - \cdots - \theta_q a_{t-q}$, the autocorrelation at lag k is calculated as:

$$\rho_k = \begin{cases} \frac{-\theta_k + \theta_1\theta_{k+1} + \cdots + \theta_{q-k}\theta_q}{1+\theta_1^2+\cdots+\theta_q^2} & k = 1, \ldots, q \\ 0 & k > q \end{cases} \tag{16}$$

3.5 The $ARMA(p, q)$ Process

3.5.1 General Definition

Combining both parts:

$$(1 - \phi_1 B - \cdots - \phi_p B^p)x_t = (1 - \theta_1 B - \cdots - \theta_q B^q)a_t \tag{17}$$

Where B is the **Backshift Operator** ($B^k x_t = x_{t-k}$).
This is often written as $\Phi(B)x_t = \Theta(B)a_t$.

3.5.2 The Partial Autocorrelation Function (PACF)

While the ACF of an AR model tails off (Gradual Decay), the **PACF** "cuts off" (Abrupt Stop)). The PACF at lag k (denoted ϕ_{kk}) is the correlation between x_t and x_{t-k} after controlling for the effects of $x_{t-1}, \ldots, x_{t-k+1}$.

- For an $AR(p)$, $\phi_{kk} = 0$ for $k > p$.
- For an $MA(q)$, the PACF tails off (opposite of the ACF).

3.6 Diagnostic Tools: ACF and PACF

To decide whether a series is an AR, an MA, or a mix of both ($ARMA$), we use two primary functions that measure how observations are related to each other over time.

3.6.1 ACF (Autocorrelation Function): The Total Correlation

The **ACF** at lag k, denoted as ρ_k, measures the simple correlation between x_t and its past value x_{t-k}. It is defined as the ratio of the autocovariance to the variance:

$$\rho_k = \frac{\gamma_k}{\gamma_0} = \frac{E[(x_t - \mu)(x_{t-k} - \mu)]}{E[(x_t - \mu)^2]} \tag{18}$$

- **Behavior for MA:** It is the "ID card" for Moving Average models. For an $MA(q)$, the ACF will be significant for q lags and then suddenly drop to zero (**cuts off**).
- **Behavior for AR:** For Autoregressive models, the ACF **tails off** (fades slowly) because the memory of the past persists through a chain reaction.
  For an $AR(1)$, this decay is geometric: $\rho_k = \phi^k$.

3.6.2 PACF (Partial Autocorrelation Function): The Direct Connection

The **PACF** at lag k, denoted as ϕ_{kk}, measures the correlation between x_t and x_{t-k} **after removing the effects** of all intervening observations ($x_{t-1}, \ldots, x_{t-k+1}$). It is formally the last coefficient of an $AR(k)$ regression:

$$x_t = \phi_{k1}x_{t-1} + \phi_{k2}x_{t-2} + \cdots + \phi_{kk}x_{t-k} + a_t \tag{19}$$

To compute ϕ_{kk} from the autocorrelations ρ_k, we use:

$$\phi_{11} = \rho_1, \quad \phi_{kk} = \frac{\rho_k - \sum_{j=1}^{k-1} \phi_{k-1,j}\rho_{k-j}}{1 - \sum_{j=1}^{k-1} \phi_{k-1,j}\rho_j} \tag{20}$$

- **Logic:** It reveals if x_{t-k} has a *direct* impact on x_t.
- **Behavior for AR:** It is the "ID card" for AR models.
  In an $AR(1)$, the PACF at lag 1 is ϕ, but for all lags $k > 1$, the PACF is 0.
  For an $AR(p)$, the PACF will be significant for p lags and then suddenly **cut off**.
- **Behavior for MA:** For Moving Average models, the PACF **tails off** (fades slowly).

3.6.3 Summary Table for Model Identification

This table is used by analysts to "read" the plots and choose the model order (p, q):

Model	**ACF Plot**	**PACF Plot**
AR(p)	Tails off (geometric decay)	**Cuts off after lag** p
MA(q)	**Cuts off after lag** q	Tails off (geometric decay)
ARMA(p,q)	Tails off	Tails off

3.7 Information Criteria: Choosing the Best Model

When we build a model, we face a trade-off: a complex model fits the current data well but may fail to forecast the future (**over-fitting**), while a simple model might miss important patterns. We use Information Criteria to find the "sweet spot" by minimizing the following:

- **AIC (Akaike Information Criterion):** $AIC = -2\ln(L) + 2k$. It provides a small penalty for each parameter (k). It is often used when the primary goal is high-quality forecasting.
- **BIC (Bayesian Information Criterion):** $BIC = -2\ln(L) + k\ln(T)$. It provides a much larger penalty that increases with the sample size (T). It is "conservative," meaning it prefers simpler models with fewer parameters to ensure the model is truly representative of the underlying process.

L **(Likelihood):** This refers to the **Maximum Likelihood**. It represents how well the model explains the observed data.

- In the criteria equations, the term $-2\ln(L)$ measures the "misfit" or unexplained information.
- As the model fit improves, L increases, causing $-2\ln(L)$ to decrease (reducing the overall "cost").

k **(Number of Parameters):** This represents the **complexity** of the model.

- For an $ARIMA(p, d, q)$ model, $k = p + q + 1$ (where $+1$ typically accounts for the constant/intercept or the noise variance).
- This serves as the **penalty term**. Adding AR or MA terms to "force" a better fit increases k, thereby increasing the penalty.

T **(Sample Size):** This is the total number of **observations** in the time series.

- In the BIC formula, the penalty is $k \ln(T)$.
- As the dataset grows (T increases), the BIC becomes significantly more "strict" than the AIC regarding the addition of new parameters.

Rule of Thumb: Calculate both for various combinations of p and q. The model with the **lowest score** is generally considered the best balance of accuracy and simplicity.

3.8 The Connection: From Identification (ACF/PACF) to Selection (AIC/BIC)

In practice, the process of building an ARMA model follows a specific logical sequence. The Information Criteria (AIC/BIC) act as the final decision-makers when the visual tools are ambiguous.

3.8.1 Step 1: The Detective Work (ACF and PACF)

We use the **ACF** and **PACF** to "shortlist" potential models:

- If the PACF cuts off at lag 2, we suspect an **AR(2)**.
- If the ACF cuts off at lag 1, we suspect an **MA(1)**.
- **The Problem:** In real-world data, the plots are rarely perfect. we might see patterns that look like they could be an $AR(2)$, an $MA(2)$, or an $ARMA(1,1)$ all at once. This is where we need a tie-breaker.

3.8.2 Step 2: The Competition (Estimation)

We estimate all the "suspect" models identified in Step 1. We now have three different versions of reality:

1. Model A: $AR(2)$ (Uses 2 parameters: ϕ_1, ϕ_2)
2. Model B: $MA(2)$ (Uses 2 parameters: θ_1, θ_2)
3. Model C: $ARMA(1,1)$ (Uses 2 parameters: ϕ_1, θ_1)

3.8.3 Step 3: The Final Judge (AIC and BIC)

Since all these models might have a similar number of parameters (k), we use **AIC** and **BIC** to see which one explains the data most efficiently.

- **Why use them with AR/MA?** If add more lags (e.g., jumping from an $AR(2)$ to an $AR(10)$), our model will always fit the training data better, but it becomes "noisy."
- **The Penalty:** AIC and BIC "punish" the model for every extra p (AR part) or q (MA part) we add.
- **Selection:** we calculate the AIC/BIC for all candidates. The one with the **lowest value** is the winner. It tells us exactly how much "memory" (AR) and how many "shocks" (MA) are truly necessary without overcomplicating the model.

3.8.4 Summary

- **PACF/ACF:** Tell us which *direction* to look in (Is it memory-based or shock-based?).
- **AIC/BIC:** Tell us exactly where to *stop* adding complexity (How many lags are enough?).

4 Part 4: ARIMA Models for Non-stationary Time Series

4.1 The Concept of Non-stationarity

Most real-world time series (like GDP, stock prices, or consumer indices) do not have a constant mean. They exhibit trends or "drifts."

4.1.1 The Random Walk (The Prototype of Non-stationarity)

A process is a **Random Walk** if the current value is the previous value plus a random shock (White Noise WN):

$$x_t = x_{t-1} + a_t, \quad a_t \sim WN(0, \sigma_a^2) \tag{21}$$

By back-substitution, assuming $x_0 = 0$:

$$x_t = a_t + a_{t-1} + a_{t-2} + \cdots + a_1 = \sum_{i=1}^{t} a_i \tag{22}$$

Properties of a Random Walk:

- **Mean:** $E[x_t] = 0$ (if no drift).
- **Variance (γ notation):** Since $Var(a_i) = \sigma_a^2$, the variance at time t is:

$$\gamma_{0;t} = Var(x_t) = t\sigma_a^2 \tag{23}$$

- **Autocovariance:** For $k > 0$, the covariance between x_t and its past x_{t-k} is:

$$\gamma_{k;t} = Cov(x_t, x_{t-k}) = (t-k)\sigma_a^2 \tag{24}$$

- **Conclusion:** Because $\gamma_{0;t}$ depends on time t (it grows to infinity), the process is **non-stationary**.

4.1.2 Theoretical Correlation and Smoothness (The Unity Property)

The autocorrelation between x_t and x_{t-k} for a random walk is given by:

$$\rho_{k,t} = \frac{\gamma_{k,t}}{\sqrt{\gamma_{0,t}\gamma_{0,t-k}}} = \frac{(t-k)\sigma_a^2}{\sqrt{(t\sigma_a^2)((t-k)\sigma_a^2)}} = \sqrt{\frac{t-k}{t}} \tag{25}$$

Crucial Observation: If t is large compared to k, then $\rho_{k,t} \approx 1$.
Implication: The sequence of x_t values will appear very **smooth** because consecutive values are almost perfectly correlated, yet the series remains **non-stationary** as the variance continues to explode ($t\sigma_a^2 \to \infty$).

4.1.3 Random Walk with Drift

If we add a constant μ, we get a **Random Walk with Drift**:

$$x_t = \mu + x_{t-1} + a_t \implies x_t = x_0 + \mu t + \sum_{i=1}^{t} a_i \tag{26}$$

This produces a **stochastic trend** where the "level" of the series changes by μ on average each period.

4.2 Integrated Processes: $I(d)$

A series x_t is said to be **integrated of order** d, denoted $x_t \sim I(d)$, if it must be differenced d times to become stationary.

4.2.1 The Difference Operator (Δ)

The first difference is defined using the Backshift Operator B:

$$\Delta x_t = (1 - B)x_t = x_t - x_{t-1} \tag{27}$$

If x_t is a Random Walk ($I(1)$), then its first difference $\Delta x_t = a_t$ is White Noise (Stationary).

4.2.2 Higher Order Differences ($I(2)$)

For some series, the "slope" itself is changing. We then use the second difference:

$$\Delta^2 x_t = \Delta(\Delta x_t) = (1 - B)^2 x_t = (1 - 2B + B^2)x_t = x_t - 2x_{t-1} + x_{t-2} \tag{28}$$

The "slope" refers to the rate of change between consecutive observations.
For example: In equation (26) The "slope" is μ.
If a series has a constant slope, it follows a straight line (a deterministic trend).
If the slope is "changing," it means the series is curving or accelerating, which usually indicates the series is Integrated of order 2, denoted as $I(2)$.

4.3 The $ARIMA(p, d, q)$ Model

The **Autoregressive Integrated Moving Average** model combines **differencing** with the ARMA structure.

4.3.1 General Equation

A process x_t is an $ARIMA(p, d, q)$, if $w_t = \Delta^d x_t$ is a **stationary** $ARMA(p, q)$ process:

$$\Phi(B)(1 - B)^d x_t = \Theta(B)a_t \tag{29}$$

Expanding the operators:

$$(1 - \phi_1 B - \cdots - \phi_p B^p)(1 - B)^d x_t = (1 - \theta_1 B - \cdots - \theta_q B^q)a_t \tag{30}$$

4.4 Determining the Order of Integration (d)

The goal of the identification stage is to find the minimum integer d such that $w_t = \Delta^d x_t$ is stationary. We follow these specific technical steps:

4.4.1 Difference between ACF and SACF

In time series analysis, it is crucial to distinguish between the theoretical properties of a process and the empirical evidence we calculate from data.

- **ACF (Autocorrelation Function):** This is a **theoretical property** of a model. It describes the population-level relationship between observations.
  - *Example:* If we **know** a process is exactly an $AR(1)$ with $\phi = 0.8$, we can use mathematical derivation to determine that $\rho_1 = 0.8$, $\rho_2 = 0.64$, and so on.

- It assumes we either have an infinite amount of data or are the "creator" of the process and know the true parameters.

- **SACF (Sample Autocorrelation Function):** This is an **empirical estimate** denoted as $\hat{\rho}_k$.
  - In reality, the true parameters (like ϕ) are unknown. we only possess a finite sample of T observations.
  - The SACF is the statistical tool used to estimate or "guess" what the underlying theoretical ACF looks like based on the available data.

4.4.2 Sample ACF (SACF)

The **SACF** ($\hat{\rho}_k$) is the empirical estimate of the autocorrelation at lag k for a series of size T. It is calculated as:

$$\hat{\rho}_k = \frac{\sum_{t=k+1}^{T}(x_t - \bar{x})(x_{t-k} - \bar{x})}{\sum_{t=1}^{T}(x_t - \bar{x})^2} \tag{31}$$

It is the primary tool for determining the order of an $MA(q)$ process (where it cuts off) and detecting non-stationarity (where it decays very slowly).

- **Non-stationary indication:** If the SACF ($\hat{\rho}_k$) decays *very slowly* or remains high for many lags, the series is non-stationary ($d \geq 1$).

- **Stationary indication:** If the SACF drops to zero exponentially or cuts off quickly, the current level of differencing is sufficient.

4.4.3 The Variance Rule

A common rule of thumb is to compare the variance of the series at different levels of differencing. We choose the d that **minimizes the variance**:

$$Var(\Delta^d x_t) < Var(\Delta^{d-1} x_t) \tag{32}$$

If differencing again *increases* the variance, we have likely "over-differenced" the series.

4.4.4 Unit Root Tests (Formal Identification)

To move beyond visual patterns, we use formal statistical tests like the **Augmented Dickey-Fuller (ADF)** test:

- **Null Hypothesis (H_0):** The series has a unit root ($\phi = 1$), meaning it is non-stationary ($d \geq 1$).

- **Alternative Hypothesis (H_1):** The series is stationary ($d = 0$).

If we fail to reject H_0, we increment d and test the differenced series Δx_t until stationarity is confirmed (p-value < 0.05).

4.4.5 Summary of Common d Values

- $d = 0$: Series is stationary around a constant mean (e.g., White Noise, ARMA).

- $d = 1$: Series has a stochastic trend (e.g., Stock prices, Random Walk).

- $d = 2$: Series has an accelerating trend or a changing slope (rare in economics, except for hyper-inflation data).

4.5 The Box-Jenkins Methodology

The course emphasizes a three-step iterative cycle for fitting ARIMA models:

1. **Identification:**
   - Determine the order of integration d (usually by looking if the ACF decays very slowly, indicating non-stationarity).
   - Use ACF/PACF of the *differenced* series to pick p and q.
2. **Estimation:** Use Maximum Likelihood Estimation (MLE) or Least Squares to find the numerical values of ϕ_i and θ_j.
3. **Diagnostic Checking:** Ensure the residuals $\hat{a}_t$ behave like **White Noise**. If there is still a pattern in the residuals, the model is incomplete. Check using the **Ljung-Box test**:

$$Q = T(T+2)\sum_{k=1}^{m} \frac{\hat{\rho}_k^2(\hat{a})}{T-k} \sim \chi^2(m-p-q) \tag{33}$$

4.6 Unit Root Testing: The Dickey-Fuller (DF) Logic

The Dickey-Fuller test is the formal statistical procedure used to determine the order of integration d. It answers the question: "Is the series a Random Walk ($I(1)$) or is it mean-reverting ($I(0)$)?"

4.6.1 1. The Basic Model and the Unit Root

Consider a simple $AR(1)$ process:

$$x_t = \phi x_{t-1} + a_t \tag{34}$$

The parameter ϕ determines the persistence of the series:

- If $|\phi| < 1$: The series is **Stationary**. Shocks eventually vanish.
- If $\phi = 1$: The series has a **Unit Root**. It is a **Random Walk**. Shocks stay in the system forever (infinite memory).

4.6.2 2. Why we rewrite the equation

In standard regression, we test if a coefficient is 0. However, testing if $\phi = 1$ directly using standard t-tests is mathematically problematic because the variance of a Random Walk is not constant.

To fix this, we subtract x_{t-1} from both sides to create a **Difference Equation**:

$$x_t - x_{t-1} = \phi x_{t-1} - x_{t-1} + a_t \tag{35}$$

$$\Delta x_t = (\phi - 1)x_{t-1} + a_t \tag{36}$$

We then define a new parameter $\gamma = (\phi - 1)$. The model becomes:

$$\Delta x_t = \gamma x_{t-1} + a_t \tag{37}$$

4.6.3 3. The Hypothesis Test

By switching to γ, we can now test for 0 instead of 1:

- **Null Hypothesis (H_0):** $\gamma = 0$ (This implies $\phi = 1$).
  $\rightarrow$ The process has a **Unit Root** (Non-stationary).
- **Alternative Hypothesis (H_1):** $\gamma < 0$ (This implies $\phi < 1$).
  $\rightarrow$ The process is **Stationary**.

4.6.4 4. The Dickey-Fuller Distribution

Crucially, under H_0, the ratio $\frac{\hat{\gamma}}{SE(\hat{\gamma})}$ (where SE is the Standard Error) does **not** follow a standard Student's t-distribution. It follows the **Dickey-Fuller Distribution**.

- If the calculated t-statistic is **more negative** than the critical value from the DF table, we reject H_0 and conclude the series is stationary ($d = 0$).
- If we fail to reject, we assume $d \geq 1$, difference the data, and potentially test again.

4.6.5 5. Augmented Dickey-Fuller (ADF)

In reality, the residuals a_t might be autocorrelated (not white noise). To fix this, we "augment" the equation by adding lags of the dependent variable (Δx_{t-j}):

$$\Delta x_t = \alpha + \beta t + \gamma x_{t-1} + \delta_1 \Delta x_{t-1} + \cdots + \delta_p \Delta x_{t-p} + a_t \tag{38}$$

This ensures that the test accounts for higher-order $AR(p)$ dynamics while still focusing on the coefficient γ to check for the unit root.

5 Part 00: Neural Networks and Deep Learning

Based on the `neuralnet.pdf` lecture, this represents a non-linear approach to regression and time series forecasting.

5.1 Foundations

- **Definition:** A neural network is a non-linear parametric regression model formed by the repeated composition of linear functions (weights W and bias b) and non-linear **Activation Functions** $\sigma(\cdot)$.
- **Activation Functions:** Common choices include **Sigmoid**, **ReLU** (Rectified Linear Unit), and **Tanh**.
- **Kolmogorov's Mapping Theorem:** States that any continuous function can be approximated by a composition of simpler functions, justifying the depth of neural architectures.

5.2 Network Architecture

- **Feed-Forward Networks:** Information flows strictly in one direction: Input $\mathbf{x} \to$ Hidden Layers $\mathbf{h} \to$ Output $\mathbf{y}$.
- **Hidden Layer Mapping:** $\mathbf{h} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$.
- **Convolutional Neural Networks (CNN):** Optimized for grid-like data using the **Convolution Operation**:

  $$[w * a]_i = \sum_k w_k a_{i-k+1} \tag{39}$$

  where w is the **kernel** (filter) that extracts spatial features.

5.3 Training and Optimization

- **Backpropagation:** The recursive application of the chain rule to compute the gradient of the loss function with respect to every weight in the network.
- **Error Decomposition:**
  - **Approximation Error:** Loss due to the restricted complexity of the chosen model class.
  - **Estimation Error:** Loss due to finite sample size and the presence of noise (leads to overfitting if too high).